

## What is Big Data?

The definition of big data is data that contains greater variety, arriving in increasing volumes and with more velocity.

## 5 Vs of Big Data

### Volume:

Huge quantity of data is generated

### Velocity:

The data is generated at a high speed

### Variety:

Different sources and types of data. Structures, unstructured and semi-structured.

### Veracity:

Truthfulness or reliability of data

### Value:

Worth of information

## What is Big Data Analytics?

It is the process of collecting, processing, cleaning, and analyzing big data to identify problems, opportunities, and increase bottom lines.

## Types of Big Data Analytics:

### Diagnostic Analysis

Understanding why something is happening

### Descriptive Analysis

Describe the current state of the business or industry

### Predictive Analysis

Make predictions about the future based on past data

### Prescriptive Analysis

Suggest solutions based on the result of descriptive or predictive analysis

## Features of Hadoop

Open Source

Network of Computers

Distributed Storage

Parallel Processing

Works on Commodity Hardware

Assumes Hardware will Fail

## Stages of Big Data Analytics

### Problem Definition:

Defining the goal, budget, and timeline of the project.

### Data Definition:

Define the data requirements and internal and external sources for getting the data.

### Data Acquisition:

Collect the data from these sources.

### Data Extraction:

Extract the most useful features of the data collected.

### Data Munging:

Converting the data into the best format.

### Data Analysis:

Analyse the data to extract meaningful pattern from the data

### Data Visualization:

Communicate the result of the analysis using visualizations

### Utilization of the Result:

Utilize the result of the analysis to drive business decisions.

## Components of Hadoop

### Hadoop Commons

Hadoop Common refers to the collection of common utilities and libraries that support other Hadoop modules.

### Hadoop Distributed File System (HDFS)

The Hadoop Distributed File System (HDFS) is a distributed file system designed to run on commodity hardware.

### MapReduce

MapReduce is a programming paradigm that enables massive scalability across hundreds or thousands of servers in a Hadoop cluster.


### Yet Another Resource Negotiator (YARN)

The fundamental idea of YARN is to split up the functionalities of resource management and job scheduling/monitoring into separate daemons.


### Hadoop Ozone

Ozone is a scalable, redundant, and distributed object store for Hadoop and Cloud-native environments.

## MapReduce



## Hadoop Distributed File System (HDFS)



## Hadoop Commands version

Returns the version of Hadoop

Syntax: `hadoop version`

### mkdir

To create a new directory on HDFS

Syntax: `hdfs dfs -mkdir <path of directory>`

### ls

To list the contents of a directory on HDFS

Syntax: `hdfs dfs -ls <path of directory>`

### put

To move content from the local file system to HDFS

Syntax: `hdfs dfs -put <path of source> <path of destination>`

### copyFromLocal

To copy data from the local file system to HDFS

Syntax: `hdfs dfs -copyFromLocal <path of source> <path of destination>`

### get

To move data from HDFS to the local file system

Syntax: `hdfs dfs -get <path of source> <path of destination>`

**copyToLocal**

To copy data from HDFS to the local file system  
 Syntax: hdfs dfs -copyToLocal <path of source> <path of destination>

**cat**

To view the contents of a file on HDFS  
 Syntax: hdfs dfs -cat <path to file>

**mv**

To move data from one location on HDFS to another.  
 Syntax: hdfs dfs -mv <path of source> <path of destination>

**cp**

To copy data from one location on HDFS to another  
 Syntax: hdfs dfs -cp <path of source> <path of destination>

**rm**

To delete data on HDFS  
 Syntax: hdfs dfs -rm -r <path to data>

**jar**

To execute MapReduce code written in a jar file  
 Syntax: hadoop jar <path of jar file> <name of class> <path of input file> <path of output file>

**Apache Pig**

Provides abstraction over MapReduce  
 Uses a scripting language called Pig Latin  
 Created by Yahoo Research in 2006  
 Moved to Apache in 2007

**Components****Parser**

Check the syntax of script and returns a Directed Acyclic Graph where the operators are represented as nodes and data flows are represented as edges

**Optimizer**


Carries out logical optimizations

**Compiler**

Compiles the optimized logical plan into a series of MapReduce jobs

**Execution Engine**

Executes the MapReduce jobs and stores the result in HDFS

**Data Model****Scalar Types**

Integers: 4-byte signed integers  
 Long: 8-byte signed integers  
 Float: Stores floating point numbers  
 Double: Double precision floating point numbers  
 Chararray: String or character array  
 Bytearray: Blob or array of bytes

**Complex Types**

Maps: Key value pairs {<key1>#<value1>, <key2>#<value2>,...}

Tuple: Collection of fields. {<field1>, <field2>, ... }

Bag: Collection of tuples. {<tuple1>, <tuple2>,...}

**Note**

Apache pig allows null values  
 If schema is available, Pig will use it for error checking and optimizations  
 If schema is not available, Pig will make the best guess based on how the script treats the data  
 Type casting syntax: (<datatype>)<variable name>

**Commands****Fs**

List all the files in HDFS  
 Syntax: fs -ls

**Clear**

Clear the interactive grunt shell  
 Syntax: clear

**History**

Syntax: history

**Reading Data**

<relation name> = LOAD <path> USING <function>  
 AS <schema>;

**Storing Data**

STORE <relation name> INTO <path> USING  
 <function>;

**Dump**

Display results on the screen for debugging  
 Syntax: dump <relation name>;

**Describe**

Helps the programmer to view the schema of the relation  
 Syntax: describe <relation name>;

**Explain**

Helps to review the logical, physical, and mapreduce execution plans  
 Syntax: explain <relation name>;

**Illustrate**

Gives step-by-step explanation of statements in Pig Execution  
 Syntax: illustrate <relation name>;

**Group**

Groups data based on the same key  
 Syntax: <grouped relation> = GROUP <relation name> BY <key>;

**Join**

Combine two or more relations  
 Syntax: <joined relation> = JOIN <relation1> BY <key1>, <relation2> BY <key2>;

**Cross**

Calculates the cross product of two or more relations  
 Syntax: <result relation> = CROSS <relation1>, <relation2>;

**Union**

Merges two relations  
 Syntax: UNION <relation1>, <relation2>;

**Filter**

Helps in filtering out the tuples out of relation, based on certain conditions

<filtered relation> = FILTER <relation name> BY <condition>;

**Distinct**

Helps in removal of redundant tuples from the relation

Syntax: <distinct relation> = DISTINCT <relation name>;

**Foreach**

Helps in generating data transformations based on column data  
 Syntax: <result relation> = FOREACH <relation name> GENERATE <keys>;

**Order by**

Displays the result in a sorted order based on one or more fields.

Syntax: <result relation> = ORDER <relation> BY <key> [DESC];

**Limit**

Gets limited number of tuples from the relation  
 Syntax: <result relation> = LIMIT <relation> <value>;

## Apache Hive

Hive is a Big Data Analytics tool that provides an SQL-like abstraction to MapReduce

### Features of Hive

Works with structured data  
Scalable and fast  
Schema is stored in a database in Hive. Processed data is stored in HDFS.  
SQL-inspired language - HiveQL  
Uses a directory structure  
Partitions and buckets enable fast retrieval

### Components of Hive

#### Hive Server

Accepts requests from Thrift, JDBC, or ODBC client

#### Beeline

Java console-based utility for connecting with relational databases and executing SQL queries

#### Hive Driver

Receives HQL statements submitted by user, creates a session, and sends the query to the compiler

#### Hive Compiler

Performs semantic analysis and type checking and then generates a directed acyclic graph

#### Hive Optimizer

Performs transformations to improve the efficiency and scalability

#### Execution Engine

Execute the plan using Hadoop

#### Metastore

Central repository which stores the metadata in the structure of tables and partitions

#### HCatalog

Table and storage management


#### WebHCat

Rest API HCatalog

## Data Model

### Tables

Data is stored in HDFS  
Metadata is stored in Hive



### Partitions

Hive organizes tables into partitions for grouping the same type of data based on a column or partition key.

The queries are faster n slices of data

### Buckets

Tables or partitions can be subdivided into buckets using functions of a column to give more structure to the data

It results in more efficient queries

Each bucket is a file in the table directory or partition directory.

## Commands

### Create

It is used to create database or table

Syntax: create database <database name>;  
create table [external] <table name>(<schema>) row format delimited fields terminated by <delimiter> stored as textfile;

### Show

Syntax: show databases;  
show tables;

### Alter

It is used to make changes to existing table

### Describe

It describes the table columns

Syntax: describe <tablename>;

### Truncate

Used to permanently truncate and delete the rows of the table

Syntax: truncate <tablename>;

### Drop

Deletes the table data

Syntax: drop <tablename>;

### Use

To select a database

Syntax: use <database name>;

### Load:

Used to move data into corresponding Hive table.

Syntax: load data [local] inpath <file path> into table <table name>;

### Insert:

Used to insert the data into Hive table

Insert Overwrite is used to overwrite the existing data in the table or partition

Insert Into is used to append the data into existing data in the table

Syntax: from <table name> insert [overwrite/into] table <table name> <query>;

### Select:

Syntax: select <columns> from <table name>;

### Grouping:

Group command is used to group the result set by one or more columns

Syntax: select <columns> from <table name> group by <column name>;

### Join

Combines fields from two tables by using values common to each

Syntax: select <columns> from <table1> <alias1> join <columns> from <table2> <alias2> on <condition>;

The result of a left outer join (or simply left join) for tables A and B always contains all records of the "left" table (A), even if the join-condition does not find any matching record in the "right" table (B).

A right outer join (or right join) closely resembles a left outer join, except with the treatment of the tables reversed. Every row from the "right" table (B) will appear in the joined table at least once.

The joined table will contain all records from both tables, and fill in NULLs for missing matches on either side.

### Quit

To exit the interactive hive shell

Syntax: quit;